
Siberian Trucking System

Software Architecture Description

Lima

athas@sigkill.dk, shantanubala@gmail.com,
jesper_tved@hotmail.com

September 29, 2012

Version history

Version	Date	Author	Comments
1	2012-08-17	KMH	Initial template based on (Rozanski and Woods, 2011)
3	2012-09-17	SHB and JTM	added chapter 3
4	2012-09-19	TRH, SHB and JTK	Revised 3.1, clarified truck transmitter concept
5	2012-09-24	TRH, SHB and JTM	added chapter 4 and 5.2

Contents

1	Introduction	4
1.1	Purpose and scope	4
1.2	Audience	5
1.3	Status	5
1.4	Architectural design approach	5
2	Glossary	6
3	System stakeholders and requirements	7
3.1	Stakeholders	7
3.2	Overview of requirements	8
3.3	System scenarios	9
4	Architectural forces	13
4.1	Goals	13
4.2	Constraints	13
4.3	Architectural principles	14
5	Architectural views	17
5.1	Context view	17
5.2	Functional view	20
5.3	Information view	22
5.4	Concurrency view	27
5.5	Deployment view	29
5.6	Development view	31
5.7	Operational view	33
6	System qualities	35
6.1	Performance and scalability	35
6.2	Security	36
6.3	Availability and resilience	36
6.4	Evolution	37
6.5	Other qualities	37

A Architecture backlog	39
B Architecture evaluation	40
C Architecture skeleton	41

Chapter 1

Introduction

Text typeset like this in an instruction on how to use the template. They should (of course) be removed in an architectural description. In the LaTeX file for the template, this can be done by renewing the `\instructions` command to output the empty string.

1.1 Purpose and scope

Russia, as the worlds largest country by land area, has an extensive raw materials industry. Since the fall of the Soviet Union, the Russian trucking industry has undergone a dramatic growth, and government initiatives aim to maintain this growth into the future. Since many of the industrial installations are located in remote areas, logistics companies face great difficulties in managing their truck fleets, particularly as transportation times are often unpredictable due to weather, accidents and the generally poor condition of the Russian road network. As an additional complication, the distant locations and low population densities prevents the use of many common communication technologies (eg. no cell phone network).

We propose the creation of a system, the Siberian Tracking System (STS), for tracking a large fleet of trucks operating in far-flung regions of Russia, owned by the (fictive) Siberian Trucking Company (STC). A central server installation will be informed of the location and state of every truck in the fleet, permitting decision making systems to have complete knowledge of the state of the company assets. Each truck is responsible for tracking its own progress, then periodically relaying it back to the main servers. The project does not involve creation of any new ground stations; trucks will use standard wireless communication methods whenever in range of appropriate networks. The transmitters on the trucks will only use one way communication to the server, with data of the trucks current position and the speed of the truck. The transmitters are supposed to be low-cost hardware,

so communication from the server to the trucks and additional measuring equipment, for instance of oil and gasoline levels have been deselected.

The STS is solely concerned with the state of the trucking fleet, and does not do freight tracking or make any kind of business logic decisions on its own. While it provides information allowing such decisions to be made, it is purely a data collection and dissemination infrastructure. In particular, it is a one-way communication system. Other means must be employed to contact the trucks on the road. Also, STS is by itself not concerned with doing data mining or presenting a sophisticated user interface to its data. Instead, a data interchange mechanism will be defined that allows other systems to receive information from STS.

1.2 Audience

The intended audience of this document, and the reason for their inclusion, are as follows.

- Business logic decision makers of STC, who must determine whether the information provided by STS is sufficient.
- Truck maintenance department representatives, to determine whether the additional equipment needed on trucks is realistic.
- The developers who have to implement the suggested architecture and design.
- Finally, whoever approves the financing of the project.

1.3 Status

The basic requirements of STS have been determined, but the actual design and implementation is not yet begun.

1.4 Architectural design approach

Explain the overall architectural approach used to describe and develop the content of the document (e.g. explain viewpoints, views and perspectives). If necessary explain the architectural views that you're using and why each is used.

Chapter 2

Glossary

Term	Definition
Mesh network	An ad-hoc (usually short-range) network formed by autonomous units within range of each other
STS	Siberian Tracking System. Refers both to the project as a whole, and to the software running on stationary servers (ie. excluding the software on truck transmitters).
STC	Siberian Tracking Company
Truck transmitter	The hardware unit on a truck. Consists of a GPS receiver and an antenna for communications.

Chapter 3

System stakeholders and requirements

3.1 Stakeholders

- **Acquirers:** the Siberian Trucking Company (STC) will be paying for the development of the system to aid in business logic. The users of the system will be members of the STC's business administration.
- **Communicators:** the technical writers who will create documentation regarding the operation of the system, while the business administration of the STC will be responsible for the training of the end users of the system.
- **Developers:** the STC has contracted a group from the University of Copenhagen to develop the system.
- **Maintainers:** the STC has a team of developers responsible for maintaining and evolving the STS system after it is completed.
- **Production Engineers:** the STS system will be deployed onto the Amazon EC2 platform, outsourcing the deployment environment to Amazon's engineering staff.
- **Suppliers:** servers will be provided by Amazon's EC2 platform, while the GPS units are assumed to already be installed on the STC's trucks.
- **Support Staff:** it is assumed that the STC has the appropriate IT staff for helping the end users of the STS in accomplishing their appropriate business administration tasks.
- **System Administrators:** Amazon provides the appropriate hardware administration while the STC has a team of developers responsible

for updating and maintaining the software environment of the EC2 instances.

- Testers: The STC has a team of developers responsible for testing and ensuring the STS works effectively.
- Users: members of the STC's business administration team who will use and analyze data provided by the STS to make informed business decisions.

3.2 Overview of requirements

Reference	Requirement description
R1	The system must provide business administrators with a detailed history of the location of every truck in the STC fleet.
R2	The system must be able to receive and store 1,000 GPS datapoints per second.
R3	The server interface for storing the trucks' location data must have an availability of at least 99 percent uptime.
R4	The tracking units present on each of the trucks must have a fallback when data cannot be sent in real-time due to bad network coverage.
R5	The server interface must be capable of receiving an individual data point (the truck's location) or a series of data points (the truck's location history over a period of time).

3.3 System scenarios

3.3.1 Functional scenarios

Scenario reference	FS1.
Overview	How truck information is sent to the server
System state	The truck is fitted with a truck transmitter and is currently driving
System environment	The system is operating normally
External stimulus	The truck transmitter has a position that should be sent to the server
Required system response	If the truck is in range of a network, the position is sent to the server and a confirmation is received. If the truck is out of range, the position is stored in the truck transmitter and will be sent when the truck is in range of a network again

Scenario reference	FS2.
Overview	How truck positions are queried by a user
System state	Positions from many different trucks have been sent to the server with timestamps
System environment	The system is operating normally
External stimulus	The user queries a specific truck, trucks within an area, and trucks that have not yet reached their designated targets on time
Required system response	The system shows the requested data as a list in the users GUI

Scenario reference	FS3.
Overview	New trucks are imported to the system
System state	New trucks are listed in the truck registry and a transmitter have been fitted into the new truck
System environment	STS, the truck transmitter and the truck registry is operating normally
External stimulus	An employee from the support staff have registered the truckID with the trucks transmitterID
Required system response	STS can now be queried for the new trucks ID

Scenario reference	FS4.
Overview	How transmission network are chosen
System state	The truck is fitted with a truck transmitter and is currently driving
System environment	The truck transmitter and the truck registry is operating normally, the truck is driving in an area without any connection
External stimulus	The truck transmitter has a new position that should be sent to the server
Required system response	The truck transmitter first tries to send via GSM mobile network, after this the long range mesh-network is tried. If neither of these worked, the position is stored to be sent at a later point.

3.3.2 System quality scenarios

Scenario reference	QS1.
Overview	If the number of trucks in the fleet increases, or if the business administrators need to collect data at a greater interval, the capabilities of the system ought to be appropriately scalable.
System environment	This process will be handled by Amazon's Elastic Load Balancing service in conjunction with Amazon EC2.
Environment changes	For example, if 10 small EC2 instances can receive and process 1,000 data points per second, the system will allocate the equivalent of 50 small EC2 instances to process an increased workload of 5,000 data points per second.
Required system behavior	When a tracking device in a truck sends data to the STS server, the data will be placed in a queue. Based on the size of the queue, an appropriate number of EC2 instances will be started or shut down to process the queue.

Software Architecture of Siberian Trucking System

Scenario reference	QS2.
Overview	If the cellular network is not available, the transmitter on the truck will fall back to using a long-distance mesh network established between the trucks.
System environment	This process will be handled by the transmitters that are installed on every truck in the STC fleet.
Environment changes	For example, if a truck leaves cell network range, it will attempt to connect to a nearby truck, which will also be connected to other trucks. As long as a single truck in the mesh has cell coverage, all of the trucks will be able to transmit their data.
Required system behavior	The trucks' transmitters will have a queue for sending data to the STS server, and will process this queue using the fastest available connection.

Scenario reference	QS3.
Overview	If part of the Amazon EC2 cluster fails, the error must not propagate to the entire system.
System environment	This robustness must be built into the software running on the EC2 cluster. Additionally, the truck transmitters must still be able to transmit data.
Environment changes	If a node in the cluster loses contact with another node, it must ensure that the overall system still has redundant data, such that further node failures can be handled. The trucks must be able to connect to another node if the usual recipient for their data fails to respond. If the system is sufficiently heavily damaged that data integrity is lost, it must still be possible for trucks to transmit data, but any queries in the data should have their result marked as incomplete.
Required system behavior	Heavy data redundancy must be built into the system. Whenever a truck reports back, it may also receive an updated list of communication endpoints.

Scenario reference	QS4.
Overview	If the transmitter on a truck fails, this failure must be noticed and rectified.
System environment	The system must have logic to detect when a truck “should” have transmitted information, but has not. We assume that the fabrication and installation of new transmitters is done externally of our system.
Environment changes	If a truck fails to report back, the maintenance department must be notified that a truck has a faulty transmitter, and that it must be replaced.
Required system behavior	We receive information from the truck register whenever the truck reaches some central locations. If the truck itself does not send the same information, its transmitter will be assumed defective, and called in for repair. As long as every truck is guaranteed to eventually stop at such a destination, any failure will eventually be discovered.

Chapter 4

Architectural forces

4.1 Goals

Business driver: Trucking in Russia is mostly done using old and inefficient trucks with bad road conditions. It is not feasible to replace the vehicles or the road system, so in order to remain competitive, the STC must optimise its logistics instead.

Project goal: The STC wishes to optimise its logistics by keeping a detailed log of truck movement. By analysing the information the log, more efficient transportation routes may be designed.

Project goal: In order to minimise truck idle time, the STC desires real-time information on truck locations, such that availability for further use can be easily predicted.

4.2 Constraints

- The company has very little in-house IT capacity, and does not wish to expand it much. In particular, it does not want to maintain its own servers.
- Since the servers must consequently be outsourced, the STS has to run in a generic, non-customised environment (such as a standard Linux server).
- The truck transmitters are very restricted, embedded hardware, that cannot run large and complicated software.
- Trucking in Russia is at an unusually high risk of hijacking. In order to not leak information about the locations of trucks, all communications has to be protected from eavesdropping, and all queries in the database must be authorised.

- As another safety-related concern, it must not be possible for a third party to falsify truck information.

4.3 Architectural principles

Principle reference	P1.
Principle statement	Redundancy of components
Rationale	The STS will be of critical importance in making real-time business decisions, so it is important that it is always available.
Implications	<ul style="list-style-type: none"> • System distributed across multiple nodes (for example, using multiple different Amazon EC2 instances across multiple availability zones). • Comprehensive error handling in all levels of the system. • Load balancing to handle node failures. • Redundancy at the data level, for example through replication.

Principle reference	P2.
Principle statement	Use of open source components
Rationale	Open source components will be used where available, in order to reduce development effort and cost.
Implications	<ul style="list-style-type: none"> • Able to modify external components, if necessary.

Principle reference	P3.
Principle statement	Encryption of communications
Rationale	To prevent outsiders from gaining knowledge about truck locations, all communications across external networks must be encrypted
Implications	<ul style="list-style-type: none"> • An encryption scheme must be decided upon. • We don't have to worry about security of the link layers.

Principle reference	P4.
Principle statement	Cryptographic signing of truck communications
Rationale	To prevent third parties from impersonating truck transmitters, all communications from trucks must be signed using assymetric cryptography.
Implications	<ul style="list-style-type: none"> • The cryptographic keys must be managed. • If a key is leaked, there must be a procedure for changing them.

Principle reference	P5.
Principle statement	The information retrieval API is REST-oriented
Rationale	REST APIs are easy to interface with existing HTTP protocol stacks.
Implications	<ul style="list-style-type: none"> • We need visible HTTP servers. • A protocol for how to represent the data in textual form must be decided.

Principle reference	P6.
Principle statement	Mesh networking by trucks
Rationale	In order to extend range beyond cell phone networks, short-range mesh networks between truck transmitters is used to indirectly transmit to STS.
Implications	<ul style="list-style-type: none">• A truck can relay information about any number of other trucks.• Every piece of truck communication must be identified by the original sender, not the truck that managed to contact the STS.

Chapter 5

Architectural views

5.1 Context view



Figure 1. Context diagram of the Siberian Transport System (STS)

5.1.1 Context diagram

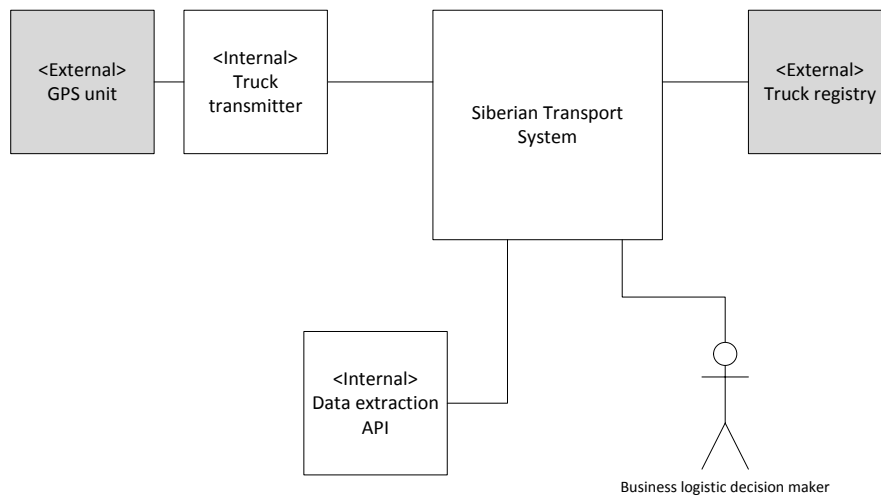


Figure 2. Context diagram with information of external and internal parts

- The truck transmitters consist of a GPS unit, that gathers the trucks coordinates, and a sender unit, which will send the coordinates to our servers through a mesh-network. The hardware is a commodity-bought external system, but the software on the sender unit is internal.
- The truck register is a system in which the Siberian trucking company's trucks are registered. All the trucks will already be registered in this system, so data is imported from here to STS. In the scope of our project, the truck register is an external system. The trucks register sends a message to STS whenever a truck is added to the fleet, removed from the fleet, or reaches one of STC's central depots.

5.1.2 Interaction scenarios

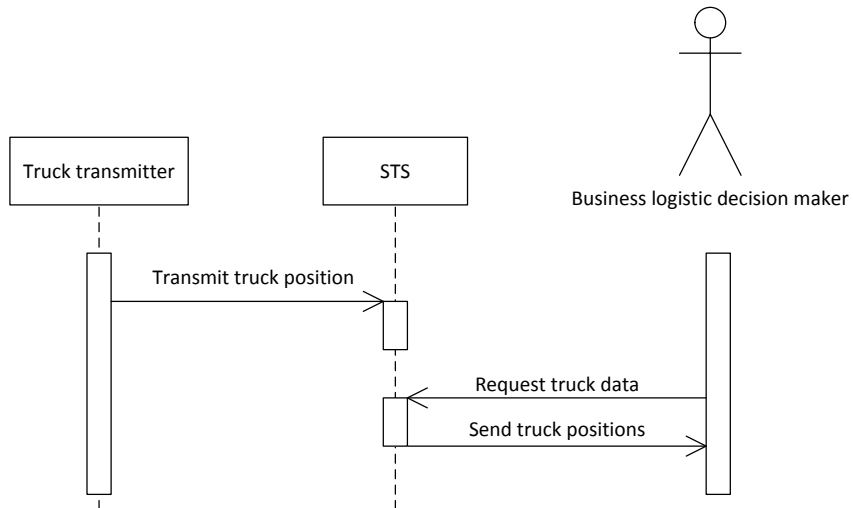


Figure 3. Truck transmitters send positions to the server. Later on, a user can make a query to see specific trucks or to see all trucks that fall behind schedule

5.2 Functional view



Figure 4. The figure shows the functional view for the truck transmitter and the Siberian Trucking System (STS)

5.2.1 Functional elements

Element name	GPS receiver
Responsibilities	acquire the current position of the truck
Interfaces – inbound	none
Interfaces – out-bound	Data collector

Element name	Data collector
Responsibilities	Gather those positions that should be sent to the server
Interfaces – inbound	GPS receiver
Interfaces – out-bound	Data sender

Software Architecture of Siberian Trucking System

Element name	Data sender
Responsibilities	send truck positions encrypted via SMS to the server and get confirmation from the data receiver
Interfaces – inbound	Data collector, Data receiver
Interfaces – out-bound	Data receiver

Element name	Data receiver
Responsibilities	Gather the positions from all the trucks, decrypt the messages and send confirmations to the trucks, when the data is received
Interfaces – inbound	Data sender
Interfaces – out-bound	Data sender and Data processing instance

Element name	Data processing instance
Responsibilities	Send data to the datastore
Interfaces – inbound	Data receiver
Interfaces – out-bound	Datastore and Truck store

Element name	Datastore
Responsibilities	Store all relevant data from the trucks
Interfaces – inbound	Data processing interface, Truck register
Interfaces – out-bound	API instance

Element name	API instance
Responsibilities	Make data accessible for other external applications
Interfaces – inbound	Datastore
Interfaces – out-bound	External client

5.2.2 Functional scenarios

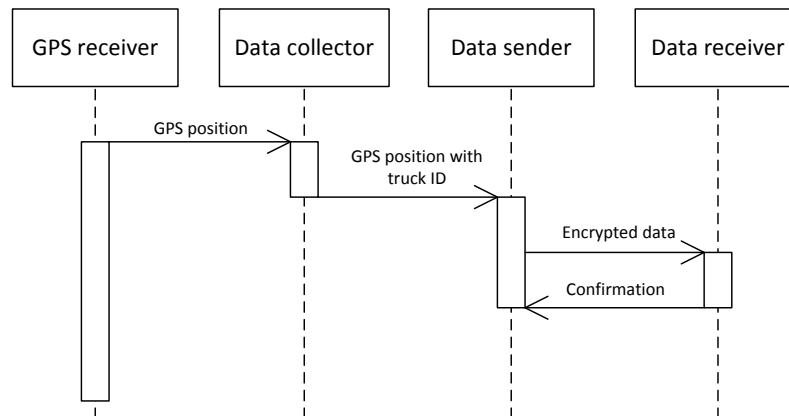


Figure 5. Truck transmitters send positions to the server end gets a confirmation when the data is received

5.2.3 System-wide processing

The transmitters will often be out of reach for some network, therefore it is essential to be able to store and send the gathered data at a later point, when the truck again is within range of a network. When the transmitter sends data, it awaits a confirmation from the server before the data is registered as sent in the truck transmitter. If a transmitter hasn't sent a position in a considerable amount of time, when the truck is expected to be at a depot within transmission range, then a notification is raised in the system, to make a technician aware of a possible problem with the trucks transmitter.

5.3 Information view

The Information view of the system defines the structure of the systems stored and transient information (e.g. databases and message schemas) and how related aspects such as information ownership, flow, currency, latency and retention will be addressed.

The GPS receiver transmits GPS coordinates along with a timestamp. This information is transmitted to the data collector, which filters and aggregates the information, as well as tagging it with a truck ID. These aggrega-

tions are sent to other truck transmitters in range, or to the STS via the data sender if possible. When the data has been successfully transmitted to the STS, it is deleted from the truck transmitter.

When truck position data is received by the STS, it becomes the property of the system. If it is to be saved (that is, if cryptographic checks affirm its validity and the truck ID is of a known truck), it will be stored in the data store component. Data is never deleted once stored.

When a truck is added or removed via the truck register, the truck ID is stored or deleted from the truck store. This data is always accessible to the data processor and must be up to date. Data never flows from the data processor to the truck store.

The user-facing API uses data from the data store and provides an aggregated view to its callers. That is, data flows from the data store to the API, never the other way. Hence, the API does not allow modification of data, it is read-only.

The data store is an unstructured NoSQL database (SimpleDB) provided by the Amazon S3 platform. The truck store is a relational database using Amazon RDS.

Every HTTP request to the API is logged by the API provider. Every component failure is logged by the redundancy infrastructure mechanism.

5.3.1 Data structure

Define or reference any architecturally significant data structures for stored and transient data, such as overview data models or message schemas. At this level you should keep the number of entities small no more than 20 or so if possible. It is not necessary to be 100% normalised for the sake of clarity it is acceptable to have some many-to-many relationships for example. Dont try and illustrate every entity and relationship here or your readers will get lost in the detail.

It may also be useful to logically group entities together that are semantically related in some way for example, all data related to customer name and address. This may help your readers to understand the data items and the relationships between them.

Here is an example data structure model which uses classic ERD notation. You can also use class diagrams here although that may be too granular a level of detail for an AD. An alternative, should you wish to use UML, is to illustrate the information structure at the package, rather than the class, level.

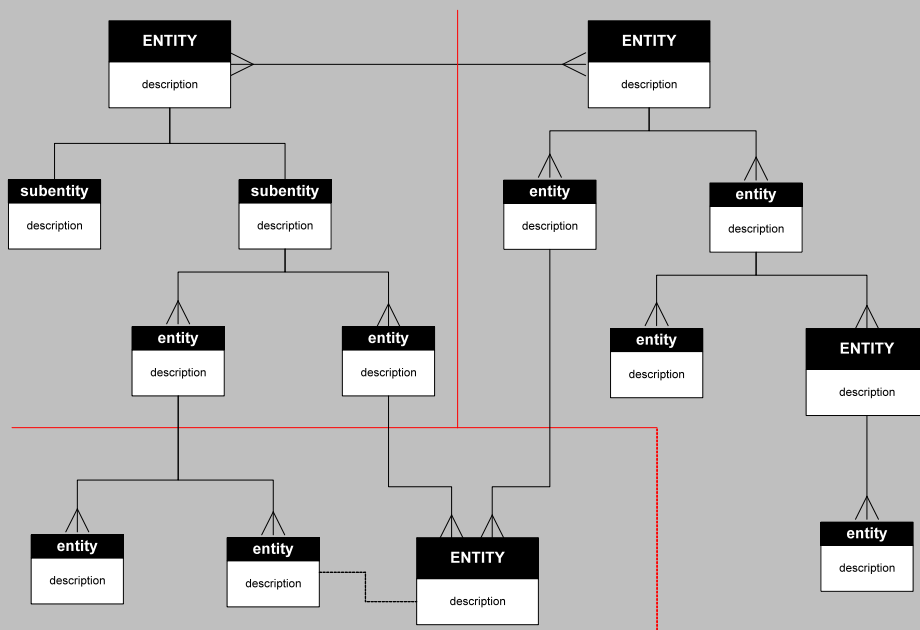


Figure 6. System data structure

The data store contains an unorganised set of tuples. Each tuple consists of a truck ID along with a timestamp and the trucks location at that time.

5.3.2 Data flow

If it is not clear from the functional views interaction diagrams, define how data flows through the system from one component to another and to external components.

As with the data structure diagram, keep this simple and focus on no more than about 10-15 key functional elements. Don't try and illustrate every data flow here or your readers will get lost in the detail.

An example is shown below using a data flow diagram.




Figure 7. System data flow

5.3.3 Data ownership

System	Truck transmitter	Data processor	API	Truck register	Data store	Truck store
Raw truck data	writer	updater	reader	none	master	none
Truck location	none	creator	reader	none	master	none
Currently known trucks	none	reader	none	writer	none	master

5.3.4 Information lifecycles

If key entities have complicated lifecycles then model the way that their state changes over time.

Focus on a few key entities whose transitions help to illuminate key features of the architecture, rather than just created / updated / updated / updated / destroyed.

There are two common techniques for modelling information lifecycles, entity life histories and state transition diagrams. Both are useful; choose one style and stick to it throughout the AD.




Figure 8. Entity life history




Figure 9. State transition

5.3.5 Timeliness and latency

If information needs to be copied around the system or is updated regularly, explain how timeliness and latency requirements will be addressed.

5.3.6 Archive and retention

Explain how will archive and retention requirements will be met by the system.

5.4 Concurrency view

The Concurrency view of the system defines the set of runtime system elements (such as operating system processes) into which the systems functional elements are packaged.

If the concurrency structure is complicated or it isn't obvious from the information in the other views, define how functional elements will be packaged into processes and threads and explain how they interact safely and reliably using suitable inter-process communication mechanisms. This can be achieved via a UML model (using stereotypes), by using a special purpose concurrency modelling language, or by creating an informal notation for the situation at hand.

Each physical truck runs a single truck transmitter instance, which contains three tasks: GPS reader, data collector and data sender.

The data and truck stores is distributed across several different processes in order to provide redundancy and scalability.

There is a 1:1 relationship between trucks and data processor tasks. Each data processor is responsible for processing data from a single truck.

The data receiver consists of a small (fixed) number of processes that merely route truck data to the proper data processors.

Each API request results in a logical task responsible for carrying out the request.

5.4.1 Concurrency model

Model the processes, process groups and threads, and the interprocess communication channels between them.

You may also choose to model the mechanisms used to protect the integrity of data and other resources shared between concurrent execution units, such as mutexes or semaphores.

You can use a UML component model to represent the information graphically, stereotyping the components appropriately.




Figure 10. Concurrency model

5.4.2 State model

Model the states that the systems runtime elements can be in, the transitions between those states and the events which drive those transitions.

A state is an identified, named stable condition which occurs during the systems runtime. An event is something that happens which causes an element to undergo a transition from one state to another. Actions may also be associated with transitions, so that while the element changes state, the action is performed.

Focus on a few key elements whose states and transitions help to illuminate key features of the architecture.




Figure 11. State model

5.5 Deployment view

The Deployment view of the system defines the important characteristics of the systems operational deployment environment. This view includes the details of the processing nodes that the system requires for its installation (i.e. its runtime platform), the software dependencies on each node (such as required libraries) and details of the underlying network that the system will require.

5.5.1 Runtime platform model

Show the systems runtime platform (defining nodes, links and the mapping of functional elements or processes to nodes).
You can use a UML deployment diagram here, or a simpler boxes-and-lines diagram.




Figure 12. Deployment model

It is often useful to explicitly map the functional elements onto the nodes that they will be running on, particularly if the deployment model is complex or the mappings arent obvious.

Functional element	Deployment node(s)

5.5.2 Software dependencies

Define the software that will be required on the various types of node in the runtime platform model, in order to support the system (such as operating system , system software or library requirements). Where versions are known you should state these.

Clearly state any known version dependencies (eg component A requires at least version X of component B).

This can usually be presented in tabular form.

5.5.3 Network model

If network requirements are complex, include a network model that illustrates the nodes, links and network hardware that the system requires, making quality of service requirements clear.

5.6 Development view

The Development view of the system defines any constraints on the software development process that are required by the architecture. This includes the systems module organisation, common processing that all modules must implement, any required standardisation of design, coding and testing and the organisation of the systems codeline.

Much of the information in this view is normally presented at a summary level, with more detail being available in other developer focused documents such as a development standards document. However you may still need to record some architecturally significant decisions at this stage, for example around choice of libraries or frameworks, or approach and tools for software deployment or configuration management.

5.6.1 Module structure

Use a model that defines the code modules that will be created and the dependencies between them. A UML package diagram is often an effective way to achieve this.




Figure 13. Module structure diagram

5.6.2 Common design

Define the common design (such as logging, security, tracing and so on) that must be performed in a standard way across the system and how it should be performed (e.g. via a design pattern or reference to a code library or sample).

5.6.3 Standards for design, code, and test

Define any standards that must be followed for design, code and unit testing, probably by reference to an external document.

5.6.4 Codeline organization

Define the codeline structure (i.e. how the source code will be held as a directory hierarchy and how it will be built into deliverable software). Define the directory hierarchy, build tools and delivery tools (such as testing or continual integration tools) that will be used to deliver the software for testing and production.

5.7 Operational view

The Operational view defines how the system will be installed into its production environment, how data and users will be migrated to it and how it will be configured, managed, monitored, controlled and supported once this is achieved. The aim of the information in this view is to show how the operational environment is to be created and maintained, rather than to define detailed instructions or procedures.

5.7.1 Installation and migration

Define the high-level steps required to install the system and any specific or unusual requirements for it.
If parallel running of old and new systems is required, explain how this will be done without disrupting existing systems, and the transition states required.

5.7.2 Operational configuration management

Define the main groups of operational configuration items and common sets of values for them (e.g. batch and overnight sets) and explain how these groups will be managed in the production environment.

5.7.3 System administration

Explain the requirements the system places on the systems administrators (in both routine and exceptional situations) and the facilities that the system will provide or rely on in the operational environment.

5.7.4 Provision of support

Define the groups involved in providing support for the system and the roles and responsibilities of each (including escalation procedures if relevant).

Chapter 6

System qualities

This section explains how the architecture presented meets its each of its required system quality properties.

While much of this information will be intrinsic to the views documented in the previous chapter, it is often useful to bring out some of it separately. In particular, if a quality property such as security or performance depends on features documented in several different views, then you should explain this here. For example, scalability may depend on optimisations in the data model (documented in the Information View) along with load balancing components (documented in the Deployment View).

6.1 Performance and scalability

For each of the main performance and scalability requirements, explain how the system will meet the requirement. Refer to practical testing and performance modelling work that has been performed as part of applying this perspective.

Requirement	How met
1. average user response time should be XX under load YY	refer to performance modelling spreadsheet

6.2 Security

For each of the main, security requirements, explain how the system will meet the requirement. Define (or reference) the threat model, security policy and security design that have been used as part of applying this perspective.

Requirement	How met
1. all users must be authenticated before being allowed to access the system	access to all screens is via standard login screen with passwords synchronised overnight to central LDAP service

6.3 Availability and resilience

Explain the A&R requirements.

Define the availability schedule(s) for the system.

Explain how the system will meet the requirements, referring to practical testing, modelling and design work that has been performed as part of applying this perspective.

Requirement	How met
1. There should be no single point of failure	all deployment nodes are clustered or load-balanced; where nodes are clustered, component failure is detected automatically and the passive node is brought up automatically

6.4 Evolution

Explain the evolution requirements.

Define the evolutionary dimensions that are relevant to the system.

Explain how the system will meet the requirements, taking into account the likelihood of each type of evolution occurring (explaining how the probabilities were arrived at) and referring to the design work performed as part of applying his perspective.

Requirement	How met
1. it must be possible to add extra input channels without having to redesign the core system	input channel components are loosely coupled to central processing modules via standardised abstract interface

6.5 Other qualities

6.5.1 Accessibility

Explain how the system meets any accessibility requirements (if any).

6.5.2 Internationalisation

Explain how the system meets any internationalisation (or localisation) requirements (if any).

6.5.3 Location

Explain how the system meets any requirements for the geographical location(s) it is to be installed in (if any).

6.5.4 Regulation

Explain how the system meets any regulatory requirements (if any).

6.5.5 Usability

Explain how the system meets any usability requirements (if any).

Appendix A

Architecture backlog

Maintain a *backlog* listing needs you have, issues and problems you have to solve, ideas for future design decisions etc. Note actions that you need to take and update the backlog as actions are completed:

Actions

- Decide on which Android version to target
- ...

Done

- Create an architectural prototype using java.nio for increased scalability

Appendix B

Architecture evaluation

Describe the architectural evaluation that you performed, cf. Chapter 14 of Rozanski and Woods (2011) including what you learned about your architecture design.

Appendix C

Architecture skeleton

Describe the architecture skeleton that you developed including how to download and run it.

Bibliography

Rozanski, N. and Woods, E. (2011). *Software Systems Architecture: Working with Stakeholders using Viewpoints and Perspectives*. Addison-Wesley Professional, second edition.